ASSIGNMENT 4.1

Task 1: Write a Map-Reduce Program to filter invalid record.

Input command: hdfs dfs –cat television.txt

Above command shows the input file television.txt

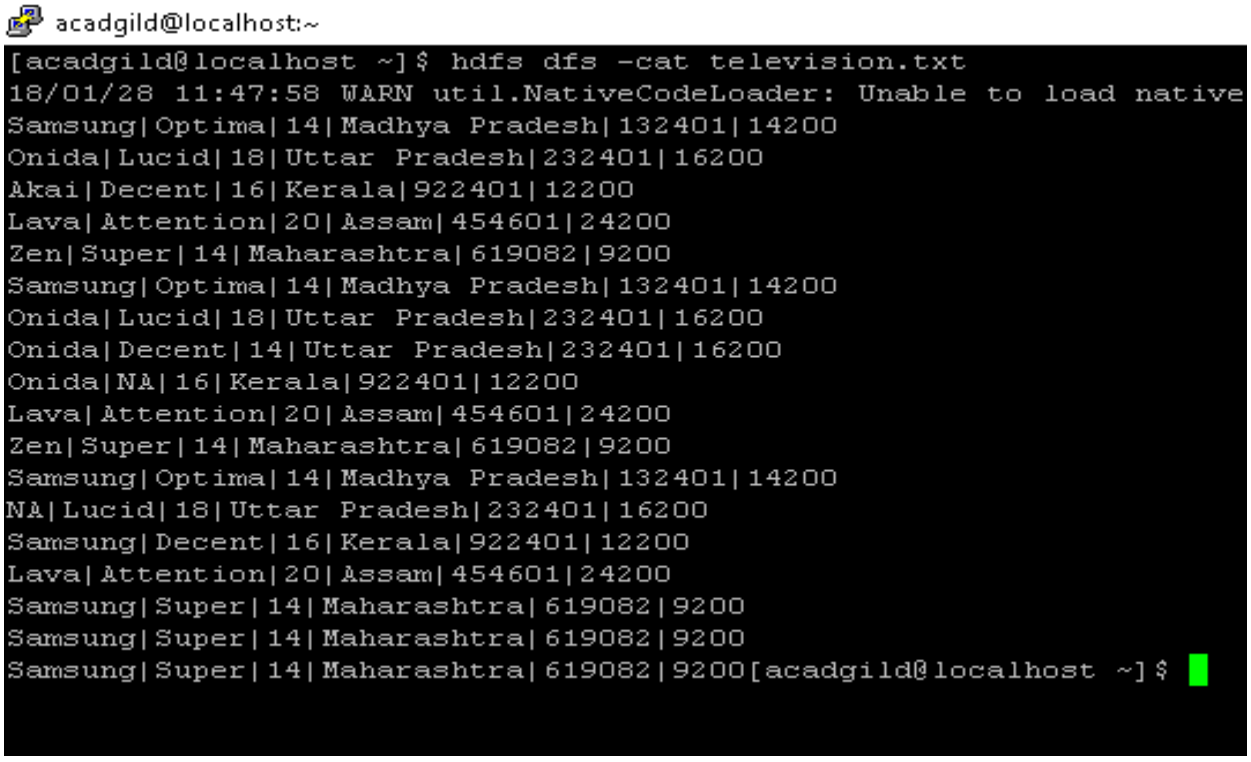

Output Command: hadoop jar /home/acadgild/Desktop/Task1.jar Task1.

InvalidRecordDriver television.txt output1

Above command is used to run the map reduce program to find invalid records. Task1.jar contains the required map-reduce program in which driver class is InvalidRecordDriver.

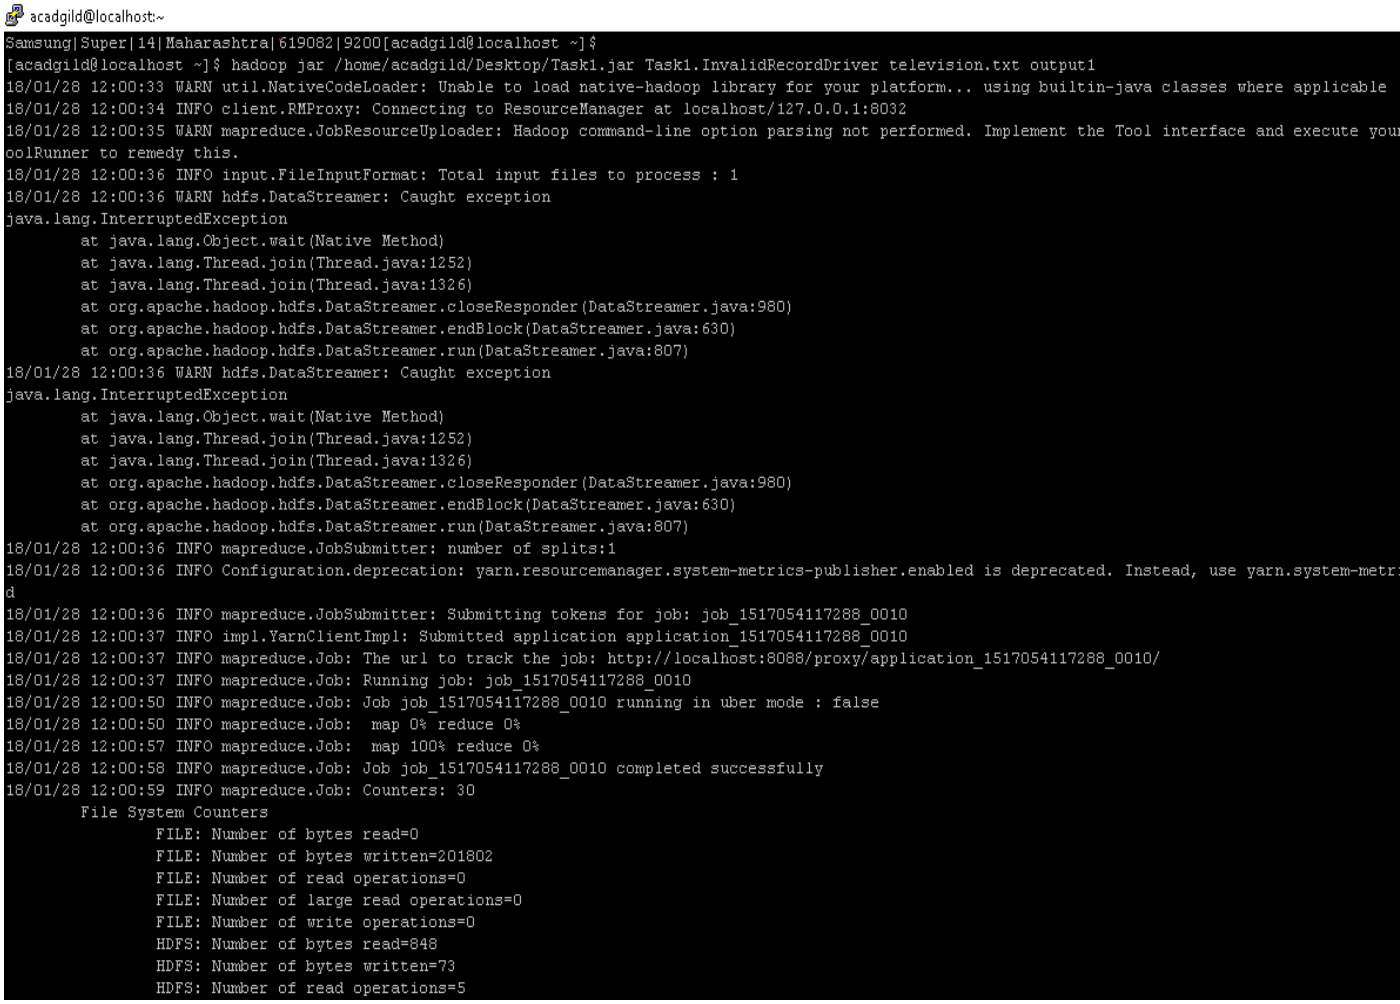

```
        File System Counters
                FILE: Number of bytes read=0
                FILE: Number of bytes written=201802
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=848
                HDFS: Number of bytes written=73
                HDFS: Number of read operations=5
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
        Job Counters
                Launched map tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=5419
                Total time spent by all reduces in occupied slots (ms)=0
                Total time spent by all map tasks (ms)=5419
                Total vcore-milliseconds taken by all map tasks=5419
                Total megabyte-milliseconds taken by all map tasks=5549056
        Map-Reduce Framework
                Map input records=18
                Map output records=2
                Input split bytes=115
                Spilled Records=0
                Failed Shuffles=0
                Merged Map outputs=0
                GC time elapsed (ms)=65
                CPU time spent (ms)=490
                Physical memory (bytes) snapshot=119857152
                Virtual memory (bytes) snapshot=2063437824
                Total committed heap usage (bytes)=62980096
        File Input Format Counters
                Bytes Read=733
        File Output Format Counters
                Bytes Written=73
```

Output: hdfs dfs –cat output1/part-m-00000

Above command will shows the output data.

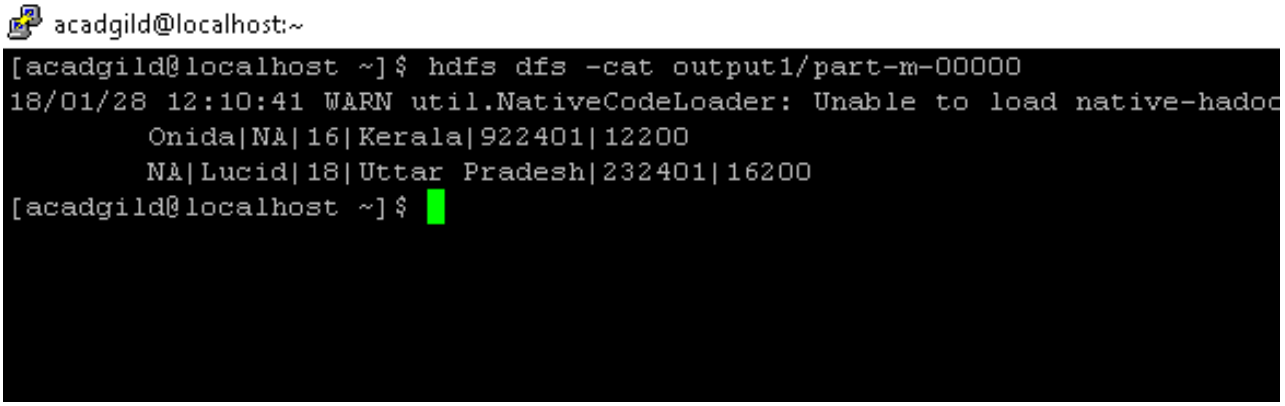

Task 2: Write a Map-Reduce Program to calculate the total unit sold for each company.

Output Command: hadoop jar /home/acadgild/Desktop/Task2.jar Task2.SoldForEachCompany_Driver television.txt output2

Above command is used to run the map reduce program to calculate total units sold for the company. Task1.jar contains the required map-reduce program in which driver class is SoldForEachCompany_Driver.

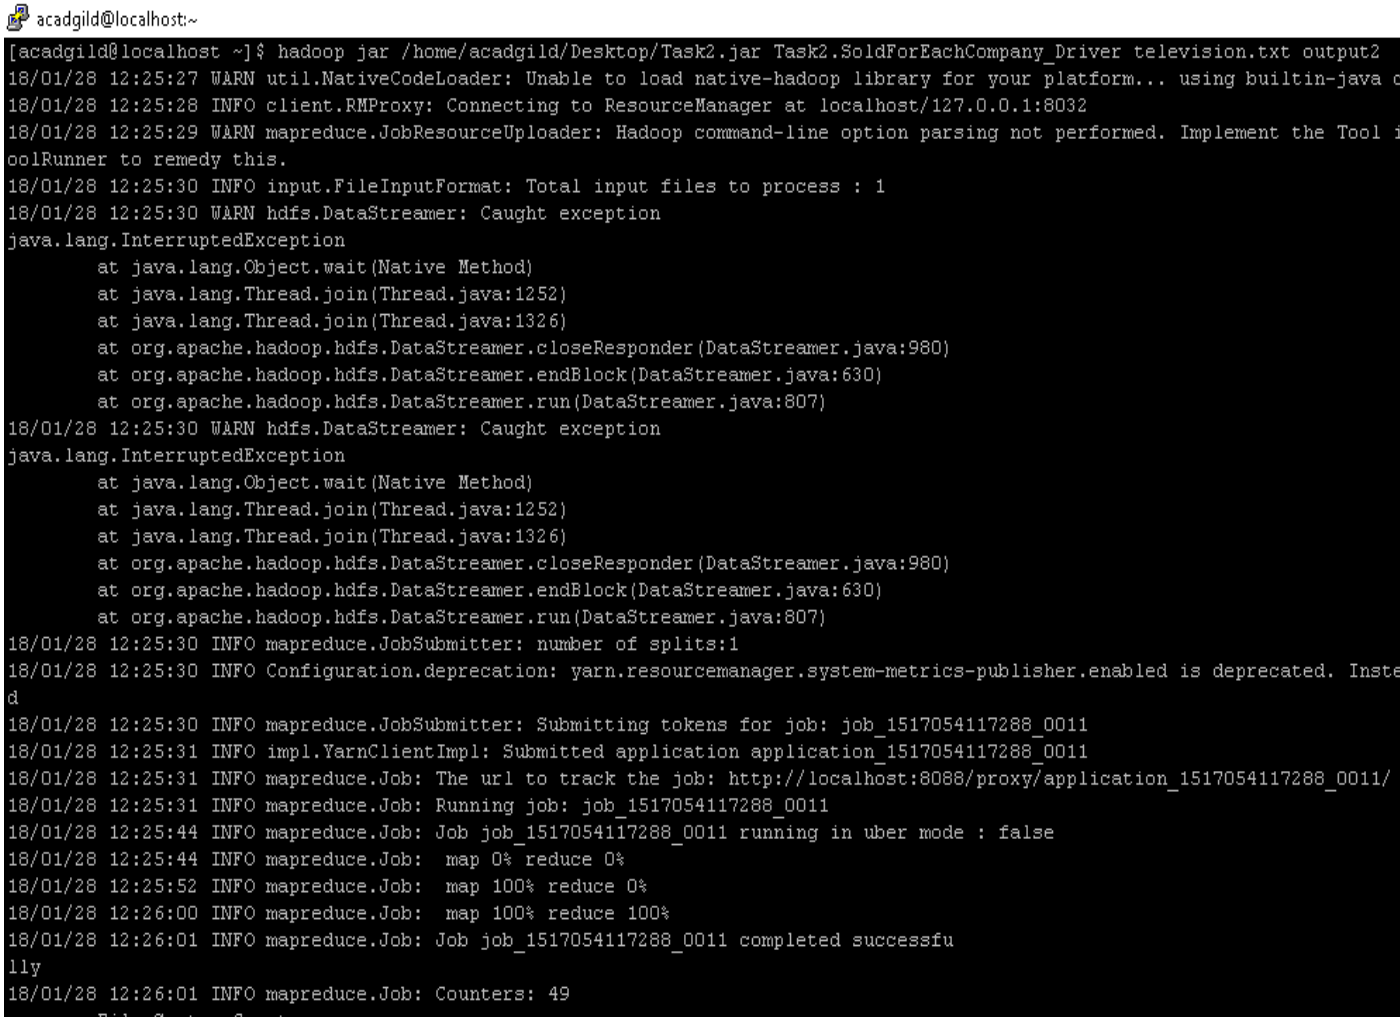

acadgild@localhost:~

```
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=5426
                Total time spent by all reduces in occupied slots (ms)=5977
                Total time spent by all map tasks (ms)=5426
                Total time spent by all reduce tasks (ms)=5977
                Total vcore-milliseconds taken by all map tasks=5426
                Total vcore-milliseconds taken by all reduce tasks=5977
                Total megabyte-milliseconds taken by all map tasks=5556224
                Total megabyte-milliseconds taken by all reduce tasks=6120448
        Map-Reduce Framework
                Map input records=18
                Map output records=18
                Map output bytes=183
                Map output materialized bytes=225
                Input split bytes=115
                Combine input records=0
                Combine output records=0
                Reduce input groups=6
                Reduce shuffle bytes=225
                Reduce input records=18
                Reduce output records=6
                Spilled Records=36
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=164
                CPU time spent (ms)=1610
                Physical memory (bytes) snapshot=348024832
                Virtual memory (bytes) snapshot=4127215616
                Total committed heap usage (bytes)=222429184
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=733
        File Output Format Counters
                Bytes Written=43
```

Output: hdfs dfs –cat output2/part-r-00000

acadgild@localhost:~

```
[acadgild@localhost ~]$ hdfs dfs -cat output2/part-r-00000
18/01/28 12:29:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
Akai    1
Lava    3
NA      1
Onida   4
Samsung 7
Zen     2
[acadgild@localhost ~]$
```

Task 3: Write a Map-Reduce Program to calculate the total unit sold in each state for Onida company.

Output Command: hadoop jar /home/acadgild/Desktop/Task3.jar Task3.SoldForOnida_Driver television.txt output3

Above command is used to run the map reduce program to calculate total units sold for onida company. Task1.jar contains the required map-reduce program in which driver class is SoldForOnida_Driver.

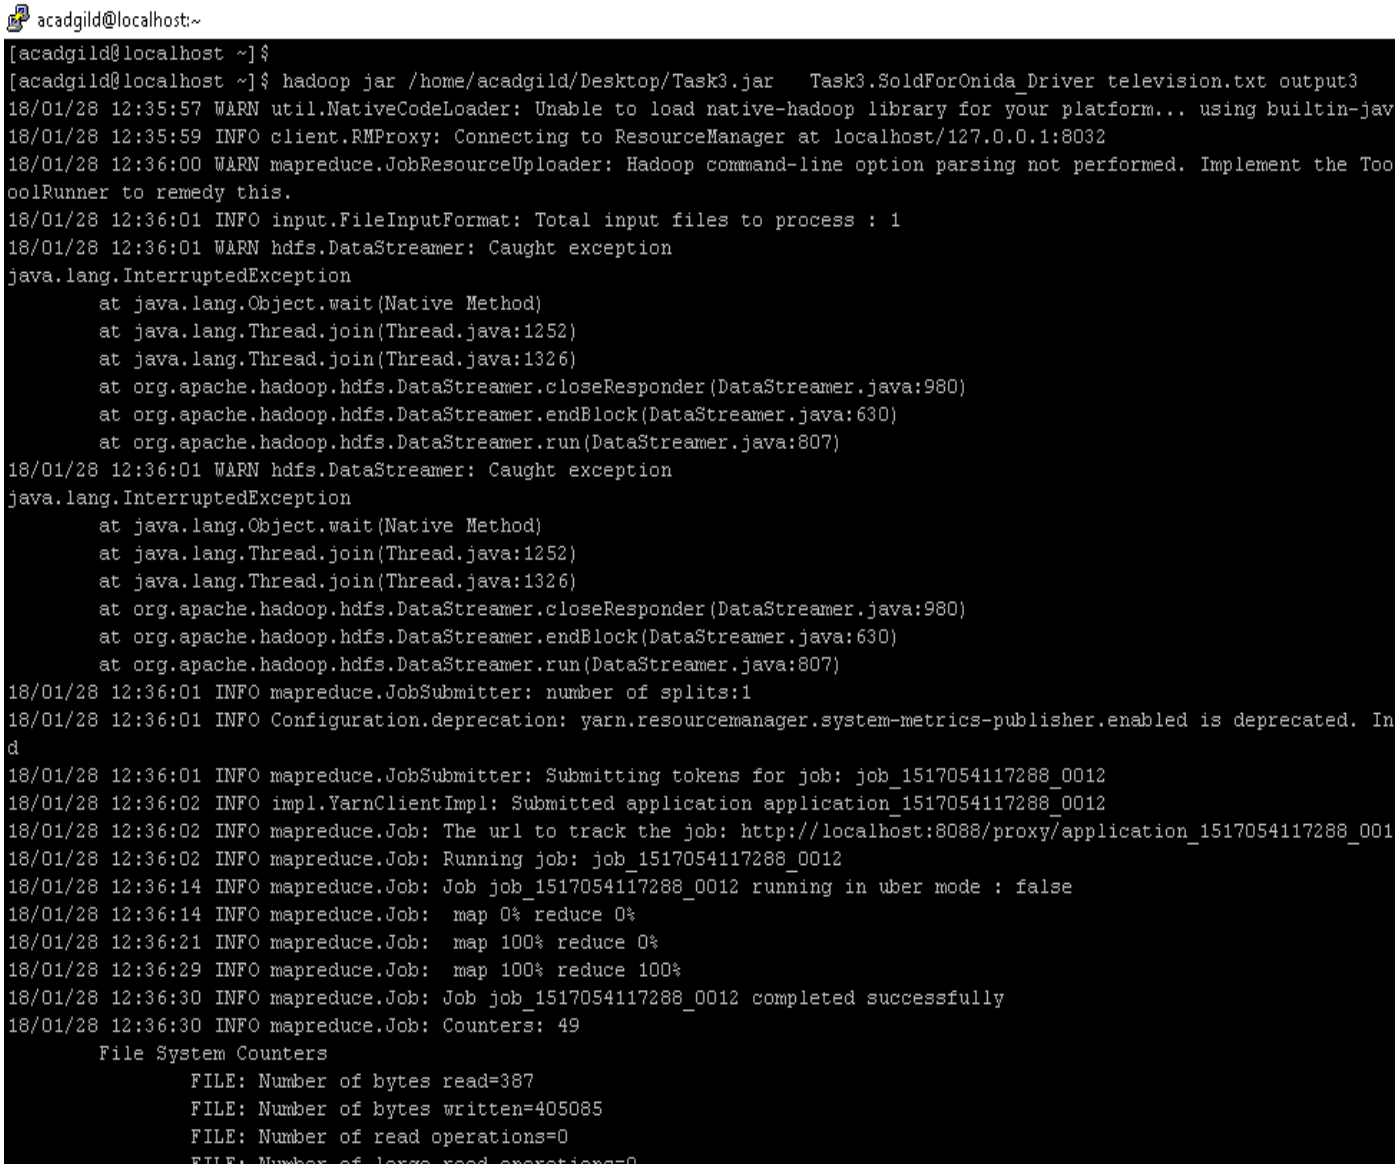

acadgild@localhost:~

```
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=5072
                Total time spent by all reduces in occupied slots (ms)=5436
                Total time spent by all map tasks (ms)=5072
                Total time spent by all reduce tasks (ms)=5436
                Total vcore-milliseconds taken by all map tasks=5072
                Total vcore-milliseconds taken by all reduce tasks=5436
                Total megabyte-milliseconds taken by all map tasks=5193728
                Total megabyte-milliseconds taken by all reduce tasks=5566464
        Map-Reduce Framework
                Map input records=18
                Map output records=22
                Map output bytes=337
                Map output materialized bytes=387
                Input split bytes=115
                Combine input records=0
                Combine output records=0
                Reduce input groups=5
                Reduce shuffle bytes=387
                Reduce input records=22
                Reduce output records=5
                Spilled Records=44
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=138
                CPU time spent (ms)=1420
                Physical memory (bytes) snapshot=348647424
                Virtual memory (bytes) snapshot=4127215616
                Total committed heap usage (bytes)=222429184
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=733
        File Output Format Counters
                Bytes Written=64
```

Output: hdfs dfs –cat output3/part-r-00000

acadgild@localhost:~

```
[acadgild@localhost ~]$ hdfs dfs -cat output3/part-r-00000
18/01/28 12:38:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... u
Assam   0
Kerala  1
Madhya Pradesh  0
Maharashtra     0
Uttar Pradesh   3
[acadgild@localhost ~]$
```